

EDS241: Assignment 3

Peter Menzies

02/19/2022

Reading in data

```
df <- read_csv(here("data", "SMOKING_EDS241.csv")) %>%
  clean_names()
```

(a) Unadjusted mean difference in birth weight of infants with smoking and nonsmoking mothers

```
mean_wt_smoking <- mean(df$birthwgt[df$tobacco == 1])
mean_wt_nonsmoking <- mean(df$birthwgt[df$tobacco == 0])

wt_diff <- mean_wt_nonsmoking - mean_wt_smoking
```

```
mdl_alc <- lm_robust(birthwgt ~ alcohol, df)
```

```
mdl_alc %>%
  tidy() %>%
  xtable()
```

```
## \begin{table}[ht]
## \centering
## \begin{tabular}{rlrrrrrrrl}
## \hline
## & term & estimate & std.error & statistic & p.value & conf.low & conf.high & df & outcome \\
## \hline
## 1 & (Intercept) & 3386.01 & 1.65 & 2056.24 & 0.00 & 3382.78 & 3389.24 & 94171.00 & birthwgt \\
## 2 & alcohol & -200.73 & 15.25 & -13.16 & 0.00 & -230.62 & -170.84 & 94171.00 & birthwgt \\
## \hline
## \end{tabular}
## \end{table}
```

```
t.test(df$alcohol[df$tobacco == 1], df$alcohol[df$tobacco == 0])
```

```
##
## Welch Two Sample t-test
```

```
##
## data:  df$alcohol[df$tobacco == 1] and df$alcohol[df$tobacco == 0]
## t = 23.838, df = 19620, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  0.03403037 0.04012795
## sample estimates:
##  mean of x  mean of y
## 0.04418246 0.00710330
```

The unadjusted mean difference in birth weight between infants with mothers who smoked during pregnancy and those with mothers who did not is 244.5393875 grams.

This would correspond to the average treatment effect of maternal smoking during pregnancy on infant birth weight under the assumption that the treatment (mother smoking during pregnancy) is statistically independent of all relevant baseline characteristics.

Evidence against this assumption: Alcohol use during pregnancy has a significant correlation with birth weight (shown above with a linear regression). The mean difference in alcohol consumption during pregnancy between smoking and non-smoking mothers in our sample is statistically different from zero (shown above with a t-test). Thus we can't make the previous assumption regarding ATE.

(b) Multiple linear regression of birth weight on smoking and all control variables

```
mdl_b <- lm_robust(birthwgt ~ tobacco + ., df)
```

```
mdl_b %>%
  tidy() %>%
  xtable()
```

| | term | estimate | std.error | statistic | p.value | conf.low | conf.high | df | outcome |
|---|-------------|----------|-----------|-----------|---------|----------|-----------|----------|----------|
| 1 | (Intercept) | 3362.26 | 12.08 | 278.41 | 0.00 | 3338.59 | 3385.93 | 94164.00 | birthwgt |
| 2 | tobacco | -228.07 | 4.28 | -53.33 | 0.00 | -236.46 | -219.69 | 94164.00 | birthwgt |
| 3 | anemia | -4.80 | 17.87 | -0.27 | 0.79 | -39.83 | 30.24 | 94164.00 | birthwgt |
| 4 | diabete | 73.23 | 13.24 | 5.53 | 0.00 | 47.29 | 99.17 | 94164.00 | birthwgt |
| 5 | alcohol | -77.35 | 14.04 | -5.51 | 0.00 | -104.87 | -49.83 | 94164.00 | birthwgt |
| 6 | mblack | -240.03 | 5.35 | -44.88 | 0.00 | -250.51 | -229.55 | 94164.00 | birthwgt |
| 7 | first | -96.94 | 3.49 | -27.79 | 0.00 | -103.78 | -90.11 | 94164.00 | birthwgt |
| 8 | mage | -0.69 | 0.37 | -1.88 | 0.06 | -1.42 | 0.03 | 94164.00 | birthwgt |
| 9 | meduc | 11.69 | 0.86 | 13.56 | 0.00 | 10.00 | 13.38 | 94164.00 | birthwgt |

The estimated coefficient on tobacco in this model is $-228.0730765 \pm 4.2767834$ grams.

(c) Estimating the effect of maternal smoking on birth weight using exact matching estimator

```
df_c <- df %>%
  mutate(mage = case_when(mage >= 34 ~ 1,
                           mage < 34 ~ 0)) %>%
  mutate(meduc = case_when(meduc >= 16 ~ 1,
                           meduc < 16 ~ 0)) %>%
  mutate(g = as.factor(paste0(mage, meduc, mblack, alcohol)))

tia <- df_c %>%
  group_by(g, tobacco) %>%
  summarise(n_obs = as.integer(n()),
            mean_wt = mean(birthwt, na.rm = T)) %>%
  gather(variables, values, n_obs:mean_wt) %>%
  mutate(variables = paste0(variables, "_", tobacco)) %>%
  pivot_wider(id_cols = g, names_from = variables, values_from = values) %>%
  ungroup() %>%
  mutate(wt_diff = mean_wt_1 - mean_wt_0,
         w_ate = (n_obs_0 + n_obs_1) / (sum(n_obs_0) + sum(n_obs_1)),
         w_att = n_obs_1 / sum(n_obs_1)) %>%
  mutate_if(is.numeric, round, 2)

xtable(tia, digits = c(0, 0, 0, 0, 2, 2, 2, 2, 2))
```

| | g | n_obs_0 | n_obs_1 | mean_wt_0 | mean_wt_1 | wt_diff | w_ate | w_att |
|----|------|---------|---------|-----------|-----------|---------|-------|-------|
| 1 | 0000 | 44274 | 13443 | 3445.69 | 3220.25 | -225.44 | 0.61 | 0.74 |
| 2 | 0001 | 214 | 448 | 3450.28 | 3124.25 | -326.03 | 0.01 | 0.02 |
| 3 | 0010 | 7007 | 1980 | 3195.97 | 3006.31 | -189.66 | 0.10 | 0.11 |
| 4 | 0011 | 71 | 226 | 3120.07 | 2817.34 | -302.73 | 0.00 | 0.01 |
| 5 | 0100 | 13425 | 535 | 3483.02 | 3273.94 | -209.08 | 0.15 | 0.03 |
| 6 | 0101 | 130 | 29 | 3510.95 | 3413.21 | -97.74 | 0.00 | 0.00 |
| 7 | 0110 | 625 | 61 | 3319.22 | 3159.05 | -160.17 | 0.01 | 0.00 |
| 8 | 0111 | 4 | 10 | 2983.50 | 3097.70 | 114.20 | 0.00 | 0.00 |
| 9 | 1000 | 5115 | 976 | 3467.41 | 3171.42 | -295.98 | 0.06 | 0.05 |
| 10 | 1001 | 56 | 45 | 3358.32 | 3097.73 | -260.59 | 0.00 | 0.00 |
| 11 | 1010 | 396 | 135 | 3185.08 | 2994.67 | -190.41 | 0.01 | 0.01 |
| 12 | 1011 | 7 | 26 | 2739.71 | 2846.38 | 106.67 | 0.00 | 0.00 |
| 13 | 1100 | 4492 | 201 | 3487.19 | 3249.45 | -237.74 | 0.05 | 0.01 |
| 14 | 1101 | 57 | 17 | 3534.91 | 3037.47 | -497.44 | 0.00 | 0.00 |
| 15 | 1110 | 147 | 19 | 3328.29 | 2852.16 | -476.13 | 0.00 | 0.00 |
| 16 | 1111 | 1 | 1 | 3459.00 | 2835.00 | -624.00 | 0.00 | 0.00 |

```
ate = sum((tia$w_ate) * (tia$wt_diff))
ate
```

```
## [1] -224.2583
```

```
att = sum((tia$w_att) * (tia$wt_diff))
att
```

```
## [1] -222.589
```

```
lm_c <- lm_robust(birthwgt ~ tobacco + g, df_c)
```

```
lm_c %>%  
  tidy() %>%  
  xtable()
```

| | term | estimate | std.error | statistic | p.value | conf.low | conf.high | df | outcome |
|----|-------------|----------|-----------|-----------|---------|----------|-----------|----------|----------|
| 1 | (Intercept) | 3445.87 | 2.23 | 1543.60 | 0.00 | 3441.50 | 3450.25 | 94156.00 | birthwgt |
| 2 | tobacco | -226.25 | 4.22 | -53.61 | 0.00 | -234.52 | -217.97 | 94156.00 | birthwgt |
| 3 | g0001 | -63.12 | 20.43 | -3.09 | 0.00 | -103.17 | -23.08 | 94156.00 | birthwgt |
| 4 | g0010 | -241.84 | 5.74 | -42.12 | 0.00 | -253.09 | -230.58 | 94156.00 | birthwgt |
| 5 | g0011 | -384.01 | 29.87 | -12.86 | 0.00 | -442.55 | -325.46 | 94156.00 | birthwgt |
| 6 | g0100 | 37.81 | 4.53 | 8.34 | 0.00 | 28.92 | 46.70 | 94156.00 | birthwgt |
| 7 | g0101 | 88.51 | 38.41 | 2.30 | 0.02 | 13.22 | 163.80 | 94156.00 | birthwgt |
| 8 | g0110 | -120.78 | 18.98 | -6.36 | 0.00 | -157.97 | -83.58 | 94156.00 | birthwgt |
| 9 | g0111 | -219.20 | 127.34 | -1.72 | 0.09 | -468.79 | 30.40 | 94156.00 | birthwgt |
| 10 | g1000 | 10.36 | 6.82 | 1.52 | 0.13 | -3.01 | 23.72 | 94156.00 | birthwgt |
| 11 | g1001 | -102.85 | 45.14 | -2.28 | 0.02 | -191.33 | -14.37 | 94156.00 | birthwgt |
| 12 | g1010 | -251.69 | 24.11 | -10.44 | 0.00 | -298.93 | -204.44 | 94156.00 | birthwgt |
| 13 | g1011 | -443.86 | 79.41 | -5.59 | 0.00 | -599.51 | -288.21 | 94156.00 | birthwgt |
| 14 | g1100 | 40.82 | 7.40 | 5.51 | 0.00 | 26.31 | 55.34 | 94156.00 | birthwgt |
| 15 | g1101 | 26.74 | 55.25 | 0.48 | 0.63 | -81.56 | 135.03 | 94156.00 | birthwgt |
| 16 | g1110 | -146.19 | 38.55 | -3.79 | 0.00 | -221.76 | -70.62 | 94156.00 | birthwgt |
| 17 | g1111 | -185.75 | 198.89 | -0.93 | 0.35 | -575.58 | 204.08 | 94156.00 | birthwgt |

The estimated average treatment effect of smoking during pregnancy on birth weight using the exact matching estimator is -224.2583 grams compared with -226.2450329 grams for the analogous linear regression.

(d) Estimating the propensity score for maternal smoking using a logit estimator

```
df_d <- df %>%
  mutate(mage2 = mage ** 2)
```

```
logit_md1 <- glm(tobacco ~ mage + mage2 + meduc + mblack + alcohol,
  family = binomial(link = "logit"),
  data = df_d)
```

```
logit_md1 %>%
  tidy() %>%
  xtable()
```

| | term | estimate | std.error | statistic | p.value |
|---|-------------|----------|-----------|-----------|---------|
| 1 | (Intercept) | 1.93 | 0.19 | 10.06 | 0.00 |
| 2 | mage | 0.08 | 0.01 | 5.21 | 0.00 |
| 3 | mage2 | -0.00 | 0.00 | -6.98 | 0.00 |
| 4 | meduc | -0.32 | 0.01 | -62.52 | 0.00 |
| 5 | mblack | -0.06 | 0.03 | -2.25 | 0.02 |
| 6 | alcohol | 2.02 | 0.06 | 33.51 | 0.00 |

```
eps <- predict(logit_md1, type = "response")
```

```
ps_wgt <- (df_d$tobacco / eps) + ((1 - df_d$tobacco) / (1 - eps))
```

(e) Using propensity score weighted least squares regression to estimate the effect of maternal smoking on birth weight

```
wls_md1 <- lm_robust(birthwgt ~ tobacco, df, weights = ps_wgt)
```

```
wls_md1 %>%
  tidy() %>%
  xtable()
```

| | term | estimate | std.error | statistic | p.value | conf.low | conf.high | df | outcome |
|---|-------------|----------|-----------|-----------|---------|----------|-----------|----------|----------|
| 1 | (Intercept) | 3425.99 | 1.85 | 1847.79 | 0.00 | 3422.36 | 3429.63 | 94171.00 | birthwgt |
| 2 | tobacco | -225.47 | 5.03 | -44.87 | 0.00 | -235.32 | -215.63 | 94171.00 | birthwgt |

The estimated effect of maternal smoking on birth weight using propensity score weighted regression is -225.4748287 grams.